

Model Optimization and Tuning Phase Template

Date	15 July 2024
Team ID	SWTID1720160264
Project Title	Predicting Compressive Strength Of Concrete Using Machine Learning
Maximum Marks	10 Marks

Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase involves refining machine learning models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

Hyperparameter Tuning Documentation (6 Marks):

Model	Tuned Hyperparameters	Optimal Values
Gradient Boosting Regression. (Randomized searchcv)	<pre>from sklearn.model_selection import RandomizedSearchCV from scipy.stats import uniform, randint gb = GradientBoostingRegressor() param_dist = { 'n_estimators': randint(100, 500), 'learning_rate': uniform(0.01, 0.2), 'max_depth': randint(1, 10), 'min_samples_split': randint(2, 20), 'min_samples_leaf': randint(1, 20), 'subsample': uniform(0.6, 0.4), 'max_features': ['auto', 'sqrt', 'log2', None] } random_search = RandomizedSearchCV(estimator=gb, param_distributions=param_dist, n_iter=100, cv=5, verbose=2, random_state=42, n_jobs=-1) random_search.fit(x_train, y_train) best_gb = random_search.best_estimator_ y_pred = best_gb.predict(x_test)</pre>	<pre>r2=r2_score(y_pred,y_test) r2</pre> <p>0.9434708311366856</p>

Gradient Boosting Regression. (Gridsearchcv)	<pre>GradientBoostingRegressor using GridSearchCV from sklearn.model_selection import GridSearchCV from sklearn.ensemble import GradientBoostingRegressor gb = GradientBoostingRegressor() param_grid = { 'n_estimators': [100, 200, 300], 'learning_rate': [0.01, 0.1, 0.05], 'max_depth': [3, 4, 5], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4] } # Initialize GridSearchCV grid_search = GridSearchCV(estimator=gb, param_grid=param_grid, cv=5, n_jobs=-1, verbose=2) # Fit the model grid_search.fit(x_train, y_train)</pre>	<pre>: best_gb = grid_search.best_estimator_ y_pred0 = best_gb.predict(x_test) : grid_search.best_score_ : 0.916712297190122</pre>
XGBoost Regression (Randomized searchcv)	<pre>XGBRegressor using RandomizedSearchCV xgb_reg = XGBRegressor() # Define the parameter distribution param_dist = { 'n_estimators': randint(100, 500), 'learning_rate': uniform(0.01, 0.2), 'max_depth': randint(3, 10), 'min_child_weight': randint(1, 10), 'subsample': uniform(0.5, 0.5), 'colsample_bytree': uniform(0.5, 0.5) } # Initialize RandomizedSearchCV random_search = RandomizedSearchCV(estimator=xgb_reg, param_distributions=param_dist, n_iter=100, cv=5, verbose=2, random_state=42, n_jobs=-1) # Fit the model random_search.fit(x_train, y_train)</pre>	<pre>: best_xgb_reg = random_search.best_estimator_ y_pred = best_xgb_reg.predict(x_test) r2 = r2_score(y_test, y_pred) r2 0.940805163311364</pre>
XGBoost Regression (Gridsearchcv)	<pre>XGBRegressor using GridSearchCV from xgboost import XGBRegressor # Initialize the XGBRegressor xgb = XGBRegressor() # Define the parameter grid param_grid = { 'n_estimators': [100, 200, 300], 'learning_rate': [0.01, 0.1, 0.05], 'max_depth': [3, 4, 5], 'min_child_weight': [1, 2, 3], 'subsample': [0.6, 0.8, 1.0], 'colsample_bytree': [0.6, 0.8, 1.0] } # Initialize GridSearchCV grid_search = GridSearchCV(estimator=xgb, param_grid=param_grid, cv=5, n_jobs=-1, verbose=2) # Fit the model grid_search.fit(x_train, y_train)</pre>	<pre>: best_xgb = grid_search.best_estimator_ y_pred = best_xgb.predict(x_test) grid_search.best_score_ 0.9225956483105566</pre>

RandomForest Regression (Randomized searchcv)	<pre>RandomForestRegressor Using RandomizedSearchCV from sklearn.ensemble import RandomForestRegressor from sklearn.model_selection import RandomizedSearchCV from scipy.stats import randint rfr = RandomForestRegressor() param_dist = { 'n_estimators': randint(100, 1000), 'max_depth': randint(1, 20), 'min_samples_split': randint(2, 20), 'min_samples_leaf': randint(1, 20), 'bootstrap': [True, False] } random_search = RandomizedSearchCV(estimator=rfr, param_distributions=param_dist, n_iter=100, cv=5, verbose=2, random_state=42, n_jobs=-1) random_search.fit(x_train, y_train)</pre>	<pre>best_rfr = random_search.best_estimator_ y_pred = best_rfr.predict(x_test) r2 = r2_score(y_test, y_pred) r2</pre> <p>0.8978707797134569</p>
RandomForest Regression (Gridsearchcv)	<pre>rfr = RandomForestRegressor() param_grid = { 'n_estimators': [100, 200, 300, 400, 500], 'max_depth': [None, 10, 20, 30, 40, 50], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4], 'bootstrap': [True, False] } grid_search = GridSearchCV(estimator=rfr, param_grid=param_grid, cv=5, n_jobs=-1, verbose=2) # Fit the model grid_search.fit(x_train, y_train)</pre>	<pre># Evaluate the model with the test set best_rfr = grid_search.best_estimator_ y_pred = best_rfr.predict(x_test) r2 = r2_score(y_test, y_pred) r2</pre> <p>0.908834984084111</p>

Performance Metrics Comparison Report (2 Marks):

Model	Baseline Metric	Optimized Metric
Gradient Boosting Regression. (Randomized searchcv)	<pre>score=r2_score(y_pred,y_test) print("R2 score: ",score*100) R2 score: 88.81108653166656 print("MSE: ",metrics.mean_squared_error(y_pred,y_test)) MSE: 27.01177924274266 print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_pred,y_test))) RMSE: 5.197285757272026</pre>	<pre>print("MAE: ",metrics.mean_absolute_error(y_pred,y_test)) print("MSE: ",metrics.mean_squared_error(y_pred,y_test)) print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_pred,y_test))) MAE: 2.646201211431639 MSE: 14.909267426711827 RMSE: 3.8612520542839244</pre>

<h3>Gradient Boosting Regression. (Gridsearchcv)</h3>	<pre>score=r2_score(y_pred,y_test) print("R2 score: ",score*100) R2 score: 88.81108653166656 print("MSE: ",metrics.mean_squared_error(y_pred,y_test)) MSE: 27.01177924274266 print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_pred,y_test))) RMSE: 5.197285757272026</pre>	<pre>print("MAE: ",metrics.mean_absolute_error(y_pred,y_test)) print("MSE: ",metrics.mean_squared_error(y_pred,y_test)) print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_pred,y_test))) MAE: 3.0454332544375693 MSE: 20.446855694652488 RMSE: 4.521819953807591</pre>
<h3>XGBoost Regression (Randomized searchcv)</h3>	<pre>print("R2 score: ",score*100) R2 score: 93.11918057409176 print("MSE: ",metrics.mean_squared_error(y_pred,y_test)) MSE: 20.446855694652488 print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_pred,y_test))) RMSE: 4.521819953807591</pre>	<pre>print("MAE: ",metrics.mean_absolute_error(y_pred,y_test)) print("MSE: ",metrics.mean_squared_error(y_pred,y_test)) print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_pred,y_test))) MAE: 2.941653629743135 MSE: 17.590176528739523 RMSE: 4.1940644402225775</pre>
<h3>XGBoost Regression (Gridsearchcv)</h3>	<pre>print("R2 score: ",score*100) R2 score: 93.11918057409176 print("MSE: ",metrics.mean_squared_error(y_pred,y_test)) MSE: 20.446855694652488 print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_pred,y_test))) RMSE: 4.521819953807591</pre>	<pre>print("MAE: ",metrics.mean_absolute_error(y_pred,y_test)) print("MSE: ",metrics.mean_squared_error(y_pred,y_test)) print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_pred,y_test))) MAE: 2.8977640539316036 MSE: 17.41153524793284 RMSE: 4.172715365897042</pre>
<h3>RandomForest Regression (Randomized searchcv)</h3>	<pre>score=r2_score(y_test,pred) print('R2 score of Random Forest Regression:',score*100) R2 score of Random Forest Regression: 91.0832388304432 print("MSE: ",metrics.mean_squared_error(pred,y_test)) MSE: 26.4968086978116595 print("RMSE: ",np.sqrt(metrics.mean_squared_error(pred,y_test))) RMSE: 5.147504927449472</pre>	<pre>print("MAE: ",metrics.mean_absolute_error(y_pred,y_test)) print("MSE: ",metrics.mean_squared_error(y_pred,y_test)) print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_pred,y_test))) MAE: 3.971962314959416 MSE: 30.34844108164759 RMSE: 5.508941920337116</pre>
<h3>RandomForest Regression (Gridsearchcv)</h3>	<pre>score=r2_score(y_test,pred) print('R2 score of Random Forest Regression:',score*100) R2 score of Random Forest Regression: 91.0832388304432 print("MSE: ",metrics.mean_squared_error(pred,y_test)) MSE: 26.4968086978116595 print("RMSE: ",np.sqrt(metrics.mean_squared_error(pred,y_test))) RMSE: 5.147504927449472</pre>	<pre>print("MAE: ",metrics.mean_absolute_error(y_pred,y_test)) print("MSE: ",metrics.mean_squared_error(y_pred,y_test)) print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_pred,y_test))) MAE: 3.637872616056214 MSE: 26.334695538705002 RMSE: 5.1317341648515855</pre>

Final Model Selection Justification (2 Marks):

Final Model	Reasoning
Gradient Boosting Regression.	In our analysis, we chose Gradient Boosting Regression due to its high accuracy and robust performance across various datasets. After hyperparameter tuning with RandomizedSearchcv, we achieved a remarkable accuracy of 94.3%, making it the best-performing model among all others we tested.